

Implementation of Quantum Approximate Optimization Algorithm for MAX2SAT Problem on Superconducting Qubits

Runzhao Guo*
(Dated: June 9, 2025)

We present a comprehensive experimental study of the Quantum Approximate Optimization Algorithm (QAOA) applied to Maximum-2-Satisfiability (MAX2SAT) problems on superconducting quantum processor. By experimenting systematically on Rigetti's Ankaa-3 quantum processor, we present how the implementation of the circuit for such an algorithm is structured, how it behaves when being scaled and how the hardware quality affects its performance.

The methodology of how the MAX2SAT problem is translated into quantum circuit is presented, including the decomposition of the problem into Hamiltonian, to universal gate and eventually into native gates of the specific quantum processor. COBYLA optimization method is explained and employed to optimize the variational parameters of the circuit.

Key findings are shown and they demonstrate that single-layer QAOA achieves a approximation ratio beyond $3/4$ for MAX2SAT. Established critical performance relationships: algorithm effectiveness varies dramatically with clause-to-variable ratio α , revealing optimal operating regimes around $\alpha \approx 1$. Scaling experiments show distinct behaviors under different constraint densities—equal variable-clause scaling leads to performance degradation from 91% to 77%, while fixed constraint scaling exhibits stabilization around 80-82% after initial drops.

Most significantly, identified a quantum advantage threshold at 75% approximation ratio, below which quantum algorithms lose computational advantage over classical random strategies. Hardware quality experiments reveal this threshold's practical importance: while high-quality 4-qubit implementations maintain 80.91% performance, low-quality conditions drop to below 75%—demonstrating the scenario where noise overwhelms quantum coherence. These results provide quantitative boundaries for near-term quantum advantage and establish benchmarks for quantum optimization algorithm viability on current hardware.

I. INTRODUCTION

Quantum computing is an rapidly evelving, interdisciplinary field of study that stand on the intersection of computer science, quantum physics, electrical engineering and mathematics. It aims to utilize properties of quantum mechanics such as superposition, entanglement, interference and perform information processing tasks that cannot be efficiently performed by classical computers.

The conceptual foundation was first laid by Paul Benioff, who proposed a quantum mechanical model of the Turing machine in 1980[1]. Richard Feynmann later proposed the idea that simulating quantum systems would require quantum mechanical information processing in 1982[2]. David Deutsch further developed the concept of universal quantum computer in 1985[3]. These early insights have laid a firm foundation for the field of quantum information science and quantum computing and lead to the first landmark quantum algorithm, Shor's algorithm, proposed by Peter Shor in 1994[4], followed by Grover's algorithm in 1996, offering quadratic speedup for unstructured search problems[5], these algorithms further demonstrated the potential of quantum computers and fueled the rapid development in the past three decades.

In the Noisy Intermediate-Scale Quantum (NISQ) era, quantum machines have made significant progress. In

the sub area of superconducting qubits, Error Correction Code such as surface code have been implemented and demonstrated to be effective in suppressing errors in quantum memory[6]. Novel Error Correction Code such as the Bivariant Bicyclic (BB) Code Have been proposed[7]. Proof of principle demonstration of BB code have been conducted on custom device[8]. Although significant progress has been made in Quantum Error Correction, the implementation of error corrected quantum algorithm remain a significant challenge, much work remains to be done in multiple areas such as coherent control, quantum gate fidelity and microwave engineering.

The near term device is however not without its importance in the context of implementation. Numerous Variational Quantum ALgorithms(VQA) have been proposed to take advantage of current quantum systems through quantum-classical hybrid optimization routine. [9–13] To date, VQAs have found use cases in various areas, including quantum chemistry simulation, machine learning and optimization. [14–18]

In particular, the Quantum Approximate Optimization Algorithm (QAOA) is one of the more popular VQAs. It is designed to find approximate solutions to hard combinatorial optimization problems on quantum computers. The algorithm encodes the problem into a related Hamiltonian, then into quantum circuit, such that the approximate solution to the problem can be constructed by projecting the quantum state onto certain operator. It leverages adiabatic time evolution and layering to optimize the variational parameters of the circuit, such that the solution to the problem can be obtained by measur-

* University of California, Los Angeles

ing the quantum state with high probability. The performance of QAOA is typically measured by the approximation ratio, which is the ratio of the cost function of the approximate solution to that of the optimal solution, denoted as $\frac{C_{QAOA}}{C_{max}}$. Theoretically, such an approximation ratio increases with the increasing number of layers, denoted as p , in the QAOA circuit, and the optimal solution can be obtained in the limit of infinite layers. However, in practice, the approximation ratio is limited by the noise in the quantum device, and the number of layers is limited by the coherence time of the quantum device.

QAOA is suitable for finding good approximate solution to several optimization problems, for instance Maximum Cut (MAXCUT)[19], Maximum Independent Set (MIS)[20] and Quadratic Unconstrained Binary Optimization (QUBO)[19], which is more general.

In this work, we focus on the Maximum 2-Satisfiability (MAX2SAT) problem, which is a special case of the MAXSAT problem. The MAX2SAT problem is NP-hard and can be formulated as a QUBO problem. The QAOA algorithm can be used to find approximate solutions to the MAX2SAT problem by encoding the problem into a related Hamiltonian and optimizing the variational parameters of the circuit.

II. BACKGROUND

In order to better present the implementation of QAOA for MAX2SAT, we first introduce the QUBO problem, the MAX2SAT problem, the classical algorithm for this problem and the QAOA algorithm for this problem.

A. Quadratic Unconstrained Binary Optimization (QUBO)

A more generalized form of the MAX2SAT problem is the Quadratic Unconstrained Binary Optimization (QUBO) problem. QUBO problem belong to NP-complete class. This mathematically assures that any NP-complete problem can be mapped to a QUBO problem in polynomial time. Research has shown that all the Karp's 21 NP-complete problems can be mapped to a QUBO counterpart[21].

In a QUBO problem, the vector of unknowns

$$\mathbf{x} = (x_1, x_2, \dots, x_n)$$

represent decision variables, or in computer science terms boolean variables, where each x_i can take the value of either 0 or 1. The problem is completed with a square symmetric matrix $\mathbf{Q} \in \mathbb{R}^{n \times n}$. The cost function is given as:

$$C(x) = \mathbf{x}^T \mathbf{Q} \mathbf{x} = \sum_{i,j=1}^n Q_{i,j} x_i x_j$$

The target of solving a QUBO problem is to find the optimal vector $\mathbf{x}^*$ such that

$$\mathbf{x}^* = \arg \min_{\mathbf{x} \in \{0,1\}^n} C(x)$$

QUBO can also be defined as maximization problem, this is done by simply negating the sign of the cost function

$$\min_{\mathbf{x} \in \{0,1\}^n} C(\mathbf{x}) = \max_{\mathbf{x} \in \{0,1\}^n} -C(\mathbf{x}) = \max_{\mathbf{x} \in \{0,1\}^n} \sum_{i,j=1}^n (-Q_{i,j}) x_i x_j.$$

It is important to note that QUBO problems are unconstrained, that is there exist no constraints on the values of the decision variables $\mathbf{x}$.

As Lucas explained in his paper [21], QUBO instances can be mapped one-to-one to Ising models. This mapping is achieved by converting the binary variables $x_i \in \{0, 1\}$ used in the QUBO formulation into spin variables $s_i \in \{-1, +1\}$ via the linear transformation $x_i = \frac{1+s_i}{2}$. When substituted into the quadratic QUBO objective function, this transformation yields a new cost function expressed entirely in terms of spin variables:

$$\sum_{i \leq j} Q_{ij} x_i x_j = \sum_{i \leq j} Q_{ij} \left(\frac{1+s_i}{2} \right) \left(\frac{1+s_j}{2} \right),$$

which expands to a quadratic form in the Ising spin variables. Grouping the resulting terms appropriately, one obtains an Ising Hamiltonian of the form

$$H(s) = - \sum_{i < j} J_{ij} s_i s_j - \sum_i h_i s_i + \text{const.},$$

where the coefficients J_{ij} and h_i are real-valued functions of the original QUBO matrix elements Q_{ij} . This transformation preserves the location of the global minimum, thereby ensuring that solving the Ising model is equivalent to solving the original QUBO problem. The resulting spin-glass representation enables direct implementation on quantum annealers and forms the basis for encoding many NP problems in the Ising model framework.

Due to this equivalence with Ising models, QUBO represents a family of problems that are suitable to be solved by quantum devices, to be more precise, quantum annealers.

As previously mentioned, many optimization problems can be mapped or reformed into QUBO problems. Although QUBO problems are limited to quadratic interactions between variables, they can be extended to higher-order terms. For example, a third-order term $x_i x_j x_k$ can be converted into a quadratic one by introducing an auxiliary variable $x' := x_i x_j$, and the original expression can be rewritten as $x' x_k$. This transformation allows for the representation of higher-order interactions in a quadratic form, making it compatible with QUBO solvers. Optimization problems with higher-order terms can be referred to as Polynomial Unconstrained Binary Optimization (PUBO) problems. Research has shown that PUBO problems can be mapped to QUBO ones efficiently[22]

B. Problem definition

Let $C = \{C_1, \dots, C_m\}$ be a weighted CNF formula in which each clause C_j carries a positive weight w_j . For a variable assignment $\sigma: \text{Var} \rightarrow \{\text{T}, \text{F}\}$ define the objective

$$f(\sigma) = \sum_{j=1}^m w_j \alpha(C_j, \sigma) \quad (1)$$

$$\alpha(C_j, \sigma) = \begin{cases} 1, & C_j \text{ is satisfied under } \sigma, \\ 0, & \text{otherwise.} \end{cases} \quad (2)$$

The MAX-SAT problem is solved by finding an assignment σ^* that maximizes the value of f . It can be proven that all MAX-SAT problem is NP-hard, that is even MAX-2-SAT is NP-hard, so to solve this problem on classical computers efficiently we should set the target to constant-factor approximations instead a entirely accurate solution.

C. Uniform random assignment (1/2-approximation)

Assign each variable independently and uniformly at random. If clause C_j contains ℓ_j different elements (in this section it is interchangeable with the expression literals), the probability it is unsatisfied equals $2^{-\ell_j}$. Hence

$$\mathbb{E}_\sigma[f(\sigma)] = \sum_j w_j (1 - 2^{-\ell_j}) \geq \frac{1}{2} \text{OPT},$$

, where OPT stands for the optimized solution value. Therefore the scheme is a $\frac{1}{2}$ -approximation. The method can be derandomised via conditional expectation without loss of guarantee, the proof of which is shown in paper published by Johnson [23]

D. Biased random assignment (≈ 0.618 -approximation)

Merge all unit clauses that involve the same element, in doing so we obtain a unique, unit clause for each variable. Then set such an element to TRUE with probability $h > 1/2$, assign its negation with probability $1 - h$, and assign all other variables uniformly. Optimising

$$\min\{h, 1 - h^2\}$$

gives $h = (\sqrt{5} - 1)/2 \approx 0.618$, which indicates an approximation to the optimal result by a factor of 0.618. This is, similar to Uniform random assignment, derandomisable. More detail on this algorithm can be found in [24].

E. LP relaxation and randomised rounding (≥ 0.632 -approximation)

Introduce fractional variables $X_i \in [0, 1]$ for every Boolean variable and $Z_j \in [0, 1]$ for every clause, by selecting according to the following simple rule

$$Z_j \leq \sum_{i \in C_j^+} X_i + \sum_{i \in C_j^-} (1 - X_i),$$

and maximise

$$\sum_j w_j Z_j$$

After solving the LP, set each variable TRUE independently with probability X_i . The number of clauses that does not satisfy the problem statement is upper bounded $\Pr[C_j \text{ unsat}] \leq (1 - Z_j/k)^k$, so the expected score is at least $(1 - 1/e) \sum_j w_j Z_j$; therefore the approximation ratio of this method is calculated as $1 - 1/e \approx 0.632$ -approximation :contentReference[oaicite:3]index=3.

F. Hybrid scheme (3/4-approximation)

Execute the LP-rounding algorithm with probability 1/2 and the uniform random assignment with probability 1/2. For every clause the expected satisfaction probability is bounded below by 3/4, giving a deterministic 3/4-approximation after derandomisation [25].

G. Relation to depth- p QAOA

At depth $p = 1$, QAOA for MAX- k -SAT produces, a distribution whose expected objective is

$$\mathbb{E}[f_{\text{QAOA}}] = (1 - 2^{-(k-1)})/2$$

for $k = 2$, i.e., 3/4 for MAX-2-SAT problem [19]. Hence the depth 1 circuit matches the best *worst-case* classical guarantee summarised above, whereas,, deeper circuits have been conjectured—but not yet proved—to beat the 3/4 barrier. Establishing such a separation therefore requires, at minimum, experimental evidence that depth- $p > 1$ QAOA outperforms all four classical baselines on *identical* problem instances.

Accounting the information presented prior in the section, a QAOA-for-MAX-SAT study should therefore include (1) measure the empirical ratio $f_{\text{alg}}/f_{\text{best}}$, or in other terms $\frac{C_{\text{QAOA}}}{C_{\text{max}}}$, on randomly generated clauses sample, and (2) report whether any quantum circuit of modest depth surpasses the 3/4 classical threshold on average or in the worst case, and analyze the performance improvement in different clauses sample configuration. (3) take into consideration the performance of the quantum circuit in the presence of noise of currently available quantum device and analyze the performance impact caused by device-specific noise configuration.




FIG. 1. Approximation ratio comparison between classical baseline algorithms and single-layer QAOA for combinatorial optimization. The approximation ratio measures the quality of solutions relative to the optimal value, with higher ratios indicating better performance. Classical baselines include **uniform random** sampling (0.50), **biased random** sampling toward better solutions (0.618), and **linear programming (LP) rounding** (0.632). The **hybrid approach** combines classical preprocessing with quantum optimization, achieving 0.75. **QAOA with p=1** matches the hybrid performance at 0.75, demonstrating that single-layer quantum circuits can compete with sophisticated classical-quantum hybrid methods. This result establishes a baseline for quantum advantage claims, showing that even shallow quantum circuits achieve approximation ratios comparable to advanced classical heuristics. The gap between simple classical methods (0.50-0.63) and quantum-enhanced approaches (0.75) illustrates the potential benefit of quantum optimization for NP-hard combinatorial problems.

H. COBYLA Optimization Method

The optimization of variational parameters in quantum algorithms like QAOA presents a fundamental challenge that distinguishes it from classical optimization problems. Unlike traditional scenarios where gradient information is readily available through analytical derivatives, quantum variational algorithms operate in environments where gradient computation is either prohibitively expensive or entirely inaccessible due to measurement noise and finite sampling statistics. This limitation necessitates the use of derivative-free optimization methods, among which the Constrained Optimization BY Linear Approximations (COBYLA) algorithm has proven particularly effective.

1. Theoretical Foundation and Algorithm Structure

COBYLA, developed by Powell in the 1990s [26], addresses the challenge of optimizing objective functions when derivative information is unavailable or unreliable. The method's core principle involves sequential linear approximation, where the optimization problem is itera-

tively simplified by constructing local linear models of both the objective function and any constraints. This transforms the original nonlinear optimization problem into a sequence of simpler linear programming subproblems.

Consider the general constrained optimization problem that COBYLA solves:

$$\text{minimize } f(\mathbf{x}) \quad (3)$$

$$\text{subject to } g_i(\mathbf{x}) \geq 0, \quad i = 1, 2, \dots, m \quad (4)$$

$$\mathbf{x} \in \mathbb{R}^n \quad (5)$$

where $f(\mathbf{x})$ represents the objective function (the expected energy from QAOA), $g_i(\mathbf{x})$ are constraint functions, and $\mathbf{x}$ represents the variational parameters γ and β .

The algorithm operates through systematic trust region management combined with linear approximation refinement. At each iteration, COBYLA maintains a set of interpolation points that serve as the foundation for constructing local linear models. The algorithm begins by establishing a trust region around the current iterate, defined by a radius parameter ρ that controls the region within which linear approximations are considered reliable.

At each iteration, COBYLA constructs linear approximations of the form:

$$f(\mathbf{x}) \approx f(\mathbf{x}_k) + \mathbf{g}_k^T(\mathbf{x} - \mathbf{x}_k) \quad (6)$$

$$g_i(\mathbf{x}) \approx g_i(\mathbf{x}_k) + \mathbf{h}_{i,k}^T(\mathbf{x} - \mathbf{x}_k) \quad (7)$$

where $\mathbf{x}_k$ is the current iterate, and $\mathbf{g}_k$ and $\mathbf{h}_{i,k}$ are estimated gradients computed using finite difference approximations. The core optimization step involves solving the linearized subproblem:

$$\text{minimize } f(\mathbf{x}_k) + \mathbf{g}_k^T(\mathbf{x} - \mathbf{x}_k) \quad (8)$$

$$\text{subject to } g_i(\mathbf{x}_k) + \mathbf{h}_{i,k}^T(\mathbf{x} - \mathbf{x}_k) \geq 0 \quad (9)$$

$$\|\mathbf{x} - \mathbf{x}_k\| \leq \rho_k \quad (10)$$

A critical component of COBYLA's effectiveness lies in its trust region management strategy. After computing a candidate step, the algorithm evaluates the actual improvement compared to the improvement predicted by the linear model through the ratio:

$$\rho_{\text{ratio}} = \frac{\text{Actual Reduction}}{\text{Predicted Reduction}} = \frac{f(\mathbf{x}_k) - f(\mathbf{x}_{k+1})}{f(\mathbf{x}_k) - m_k(\mathbf{x}_{k+1})}$$

When this ratio is close to 1, the linear approximation accurately captures local behavior, justifying an increased trust region radius. When the ratio is significantly less than 1, it triggers a reduction in radius to enforce more conservative steps.

2. Application to Quantum Optimization

In QAOA parameter optimization, COBYLA offers several distinct advantages. The derivative-free nature

eliminates the need for parameter-shift rules or finite difference approximations to compute gradients, which can be computationally expensive and numerically unstable with measurement noise. The algorithm's constraint handling capabilities prove valuable for incorporating parameter bounds that arise naturally in quantum algorithms.

The algorithm's robustness to noise is particularly important in near-term quantum devices, where measurement statistics are finite and gate fidelities are imperfect. COBYLA's reliance on function value comparisons rather than precise gradient information makes it inherently more tolerant of fluctuations characteristic of real quantum hardware implementations.

From a computational perspective, COBYLA typically requires $O(n^2)$ function evaluations per iteration, where n is the parameter space dimension. While this scaling might appear unfavorable compared to gradient-based methods requiring $O(n)$ evaluations per iteration, the absence of gradient computation costs often makes COBYLA competitive in practice for quantum applications. The algorithm's memory requirements are modest, scaling as $O(n^2)$, making it practical even for high-dimensional parameter spaces encountered in deep QAOA circuits.

Convergence rates for COBYLA are typically sublinear, with theoretical guarantees providing convergence to first-order critical points under standard regularity assumptions. While slower than superlinear convergence achievable by second-order methods, practical performance often compares favorably due to reduced per-iteration cost and robustness properties. The effectiveness depends significantly on proper parameter tuning, particularly the initial trust region radius and tolerances, which should be chosen considering the expected magnitude of quantum measurement noise and the typical scale of optimization variables.

The implementation of COBYLA in this work is done by importing the minimize function from the scipy optimization library, which include the COBYLA algorithm as one of the optimization methods.

III. QAOA FOR MAX2SAT

In this section we present how the circuit for QAOA for MAX2SAT is constructed in universal gate set, how it is decomposed into native gates for the specific processor we are using, the optimization of the circuit as well as measurement is also covered.

A. QAOA Circuit Construction for MAX2SAT Problems

This subsection presents the systematic construction of quantum circuits implementing the Quantum Approximate Optimization Algorithm (QAOA) for Maximum 2-

Satisfiability (MAX2SAT) problems. The exposition focuses on the practical translation from boolean constraint satisfaction to executable quantum operations, emphasizing the underlying mathematical structure that enables efficient quantum implementation.

1. Problem Encoding and Mathematical Foundation

The construction of QAOA circuits for MAX2SAT begins with the systematic translation of boolean satisfiability constraints into quantum operators. Consider a MAX2SAT instance with n variables $\{x_1, x_2, \dots, x_n\}$ and m clauses, where each clause contains exactly two literals. The classical objective function counts the number of satisfied clauses and seeks to maximize this quantity.

The quantum encoding maps each boolean variable x_i to a qubit, creating an n -qubit system where computational basis states $|z_1 z_2 \dots z_n\rangle$ correspond to truth assignments. The key insight lies in expressing each clause as a penalty term that contributes zero energy when satisfied and positive energy when violated.

For a clause $(\neg x_i \vee x_j)$, the corresponding quantum operator becomes

$$\frac{1}{4}(I - Z_i)(I + Z_j), \quad (11)$$

where Z_i represents the Pauli-Z operator acting on qubit i . This construction ensures that the term evaluates to zero when the clause is satisfied (either $x_i = 0$ or $x_j = 1$) and to $\frac{1}{2}$ when violated. The complete cost Hamiltonian emerges as:

$$H_C = \sum_{k=1}^m \frac{1}{4}(I - Z_{i_k})(I + Z_{j_k}), \quad (12)$$

where the sum extends over all clauses k , and (i_k, j_k) denotes the variable indices in clause k . Expanding this expression yields a sum of local Z and ZZ terms that can be efficiently implemented on quantum hardware.

B. Circuit Architecture and Decomposition

The QAOA circuit consists of p alternating layers of cost and mixing unitaries, applied to an initial uniform superposition state. The cost unitary $U_C(\gamma) = \exp(-i\gamma H_C)$ encodes the problem structure, while the mixing unitary $U_M(\beta) = \exp(-i\beta H_M)$ with $H_M = \sum_i X_i$ enables transitions between computational basis states.

The decomposition of $U_C(\gamma)$ requires careful attention to the commutation properties of the constituent terms. Since terms acting on disjoint qubits commute, the cost unitary factors as:

$$U_C(\gamma) = \prod_{k=1}^m \exp\left(-i\frac{\gamma}{4}(I - Z_{i_k})(I + Z_{j_k})\right). \quad (13)$$

Each individual term further decomposes using the identity

$$\exp\left(-i\frac{\theta}{4}(I-Z)(I+Z')\right) = \exp\left(-i\frac{\theta}{4}\right) \exp\left(-i\frac{\theta Z}{4}\right) \times \exp\left(i\frac{\theta Z'}{4}\right) \exp\left(i\frac{\theta ZZ'}{4}\right), \quad (14)$$

yielding a sequence of single-qubit rotations and controlled operations.

The practical implementation proceeds by expanding $(I - Z_i)(I + Z_j) = I + Z_j - Z_i - Z_i Z_j$, leading to the unitary:

$$\exp\left(-i\frac{\gamma}{4}(I + Z_j - Z_i - Z_i Z_j)\right) = \exp\left(-i\frac{\gamma}{4}\right) \exp\left(-i\frac{\gamma Z_i}{4}\right) \times \exp\left(i\frac{\gamma Z_i}{4}\right) \exp\left(i\frac{\gamma Z_i Z_j}{4}\right). \quad (15)$$

The global phase factor $\exp(-i\gamma/4)$ can be omitted, leaving three nontrivial rotations that translate directly to quantum gates.

C. Gate Sequence Implementation

The translation to elementary quantum gates follows established decomposition patterns. Single-qubit rotations $\exp(\pm i\gamma Z_i/4)$ implement as $R_Z(\pm\gamma/2)$ gates, while two-qubit terms $\exp(i\gamma Z_i Z_j/4)$ require CNOT-based constructions.

The standard decomposition of $\exp(i\theta Z_i Z_j)$ employs the identity $Z_i Z_j = Z_i \otimes Z_j$, leading to the gate sequence: **ZZ Rotation Decomposition:**

1. Apply CNOT(i, j)
2. Apply $R_Z(\theta)$ on qubit j
3. Apply CNOT(i, j)

This construction leverages the fact that CNOT gates transform the ZZ interaction into a local Z rotation on the target qubit when the control is in eigenstate $|1\rangle$.

For each clause $(\neg x_i \vee x_j)$, the complete gate sequence within the cost layer becomes:

Single Clause Implementation:

1. Apply $R_Z(-\gamma/2)$ on qubit i
2. Apply $R_Z(\gamma/2)$ on qubit j
3. Apply CNOT(i, j)
4. Apply $R_Z(\gamma/2)$ on qubit j
5. Apply CNOT(i, j)

The mixing layer implementation proves more straightforward, consisting of $R_X(\beta)$ rotations applied simultaneously to all qubits, since all X operators mutually commute.

D. Circuit Optimization and Depth Reduction

Real quantum implementations benefit significantly from circuit optimization techniques that reduce gate count and depth while preserving the mathematical structure. The commutation properties of non-overlapping clause operators enable parallel gate application, substantially reducing circuit depth for sparse constraint graphs.

When multiple clauses share common variables, careful scheduling of operations can merge consecutive single-qubit rotations. For instance, if variable x_i appears in multiple clauses with coefficients $\{\alpha_k\}$, the individual R_Z rotations can be combined into a single $R_Z(\sum_k \alpha_k)$ operation. The two-qubit ZZ rotations present greater optimization challenges due to their CNOT gate requirements. Advanced compilation techniques exploit hardware connectivity constraints and gate error rates to determine optimal qubit routing and gate scheduling. Native gate sets of specific quantum devices may offer alternative decompositions that reduce implementation overhead.

E. Concrete Circuit Example

Consider the simple MAX2SAT instance with three variables and two clauses: $C_1 = (\neg x_1 \vee x_2)$ and $C_2 = (x_1 \vee \neg x_3)$. The cost Hamiltonian becomes:

$$H_C = \frac{1}{4}[(I - Z_1)(I + Z_2) + (I + Z_1)(I - Z_3)]. \quad (16)$$

Expanding yields:

$$H_C = \frac{1}{2}I + \frac{1}{4}(Z_2 - Z_1 - Z_1 Z_2 + Z_1 - Z_3 - Z_1 Z_3). \quad (17)$$

The corresponding cost unitary circuit for parameter γ consists of:

a. *Clause C_1 contribution:*

- $R_Z(\gamma/2)$ on qubit 1
- $R_Z(-\gamma/2)$ on qubit 2
- CNOT(1, 2), $R_Z(-\gamma/2)$ on qubit 2, CNOT(1, 2)

b. *Clause C_2 contribution:*

- $R_Z(-\gamma/2)$ on qubit 1
- $R_Z(\gamma/2)$ on qubit 3
- CNOT(1, 3), $R_Z(-\gamma/2)$ on qubit 3, CNOT(1, 3)

The single-qubit rotations on qubit 1 from both clauses combine into $R_Z(\gamma/2 - \gamma/2) = R_Z(0)$, effectively canceling out. This demonstrates how optimization can eliminate redundant operations in practical implementations.

F. Parameterization and Measurement

The complete p -layer QAOA circuit takes the form:

$$|\psi(\gamma, \beta)\rangle = U_M(\beta_p)U_C(\gamma_p)\cdots U_M(\beta_1)U_C(\gamma_1)|+\rangle^{\otimes n}, \quad (18)$$

where $|+\rangle^{\otimes n}$ represents the uniform superposition initial state created by applying Hadamard gates to $|0\rangle^{\otimes n}$. The $2p$ parameters $\{\gamma_1, \dots, \gamma_p, \beta_1, \dots, \beta_p\}$ undergo classical optimization to maximize the expectation value $\langle\psi(\gamma, \beta)|H_C|\psi(\gamma, \beta)\rangle$.

Measurement proceeds by sampling the computational basis states and evaluating the classical objective function for each outcome. The measurement statistics provide an unbiased estimator of the expectation value, enabling gradient-based optimization of the variational parameters.

Modern quantum devices implement this measurement through repeated circuit execution, accumulating statistics across many shots. The sampling overhead scales with the desired precision, typically requiring thousands of measurements for accurate parameter optimization.

The circuit construction presented here provides a systematic pathway from MAX2SAT problem instances to executable quantum circuits, enabling practical experimentation with QAOA algorithms on near-term quantum devices. The modular structure facilitates both classical simulation and quantum hardware implementation, supporting ongoing research into quantum optimization advantages.

IV. CIRCUIT DIAGRAM AND SIMULATION

Now that a clear map from MAX2SAT problem to executable quantum circuit has been established, we can proceed to a proof-of-principle simulation of this QAOA on MAX2SAT problem.

A. Circuit Diagram

The complete circuit diagram for QAOA for MAX2SAT problem can be decomposed into three parts, the first part is the initial state preparation, which prepares the uniform superposition state $|+\rangle^{\otimes n}$. This section is implemented by applying Hadamard gates to all qubits. In the Rigetti Ankaa-3 quantum processor, single qubit operations are implemented by applying R_x , R_y rotations via microwave pulses and R_z rotations conducted virtually by applying phase shifts to the following operations in compilation process. Therefore the Hadamard gate can be decomposed into three pulses, $R_y(\pi/4)$ - $R_x(\pi/2)$ - $R_y(-\pi/4)$. In favor of convenience and clarity, we will use the H symbol to denote this three-pulse sequence in the following discussion.

The second part of the circuit is the cost unitary or the separator, which in theory is done by implementing

controlled-Rotation gates. In practice, the Ankaa-3 quantum processor supports iSWAP gate and can be used to implement CNOT gates with the help of Rotation gates. The controlled-Rotation gates are then decomposed into a sequence of CNOT gates and R_z gates (which does not introduce additional pulse on the physical qubit and is therefore disireable).

The third part of the circuit is the mixing unitary, which is essentially a sequence of rotations on the x - y plane (or transverse plane for generalization). In this work we explore the cases where the mixing unitary is implemented by applying R_x rotations to all qubits.

As mentioned in the previous section, the cost and mixing unitaries are parameterized by γ and β respectively, the optimization of these parameters is done by classical optimization algorithms.

B. Single Qubit Open System Simulation Framework

The simulation of single qubit open quantum systems requires a comprehensive framework that captures both the coherent evolution under a Hamiltonian and the incoherent effects of environmental interactions. Our approach combines the mathematical formalism of Kraus operators with numerical integration of the master equation to provide a complete picture of quantum decoherence processes.

1. Quantum State Representation and Operators

We begin by establishing the fundamental building blocks of our single qubit system. The computational basis states are represented as density matrices:

$$|0\rangle\langle 0| = \begin{pmatrix} 1 & 0 \\ 0 & 0 \end{pmatrix}, \quad |1\rangle\langle 1| = \begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix} \quad (19)$$

The Pauli matrices serve as the fundamental operators for describing single qubit dynamics:

$$\sigma_x = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad \sigma_y = \begin{pmatrix} 0 & -i \\ i & 0 \end{pmatrix}, \quad \sigma_z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix} \quad (20)$$

These operators form a complete basis for describing any single qubit operation and measurement, making them essential tools for constructing both the system Hamiltonian and the noise processes.

2. System Hamiltonian Construction

The coherent evolution of our qubit is governed by a Hamiltonian that combines energy splitting and driving terms:

$$H = -\frac{\delta}{2}\sigma_z + \frac{\omega}{2}\sigma_x \quad (21)$$

This Hamiltonian captures two physical effects: the σ_z term represents the energy difference between the qubit's computational basis states (characterized by the detuning parameter δ), while the σ_x term describes coherent driving or tunneling between states (with strength ω). This form is particularly relevant for systems like superconducting qubits or trapped ions, where both energy splitting and coherent control are present.

3. Environmental Noise Models via Kraus Operators

Environmental interactions are modeled using the Kraus operator formalism, which provides a complete description of how quantum information is lost to the environment. Each noise channel is characterized by a set of Kraus operators $\{K_i\}$ that satisfy the completeness relation $\sum_i K_i^\dagger K_i = I$, ensuring that the evolution preserves the trace of the density matrix.

Our simulation framework incorporates four fundamental decoherence mechanisms:

Bit Flip Channel: This channel models random flips of the qubit state, analogous to classical bit errors. The Kraus operators are:

$$K_0 = \sqrt{1-p} I, \quad K_1 = \sqrt{p} \sigma_x \quad (22)$$

where p represents the probability of a bit flip occurring.

Amplitude Damping: This channel describes energy relaxation, where the excited state $|1\rangle$ spontaneously decays to the ground state $|0\rangle$. The Kraus operators take the form:

$$K_0 = \begin{pmatrix} 1 & 0 \\ 0 & \sqrt{1-p} \end{pmatrix}, \quad K_1 = \begin{pmatrix} 0 & \sqrt{p} \\ 0 & 0 \end{pmatrix} \quad (23)$$

Phase Damping: This channel models pure dephasing, where the relative phase between computational basis states is randomly altered without affecting populations:

$$K_0 = \sqrt{1-p} I, \quad K_1 = \sqrt{p} \sigma_z \quad (24)$$

Depolarizing Channel: This represents the most general form of decoherence, where the qubit randomly evolves toward the maximally mixed state:

$$K_0 = \sqrt{1-\frac{3p}{4}} I, \quad K_1 = \sqrt{\frac{p}{4}} \sigma_x, \quad (25)$$

$$K_2 = \sqrt{\frac{p}{4}} \sigma_y, \quad K_3 = \sqrt{\frac{p}{4}} \sigma_z \quad (26)$$

4. Numerical Integration and Master Equation Approach

The time evolution of our open quantum system is computed by solving the master equation, which describes how the density matrix evolves under the combined influence of coherent Hamiltonian dynamics and

incoherent environmental effects. Rather than applying Kraus operators at discrete time steps, we employ a continuous-time approach using the Lindblad master equation:

$$\frac{d\rho}{dt} = -i[H, \rho] + \sum_k \gamma_k \left(L_k \rho L_k^\dagger - \frac{1}{2} \{L_k^\dagger L_k, \rho\} \right) \quad (27)$$

Here, L_k are the Lindblad operators derived from our Kraus operators, and γ_k are the corresponding decay rates. This approach provides a more physically realistic description of continuous environmental monitoring compared to discrete noise applications.

The numerical integration is performed using adaptive timestep methods that ensure both accuracy and computational efficiency. By leveraging specialized quantum optics libraries, we can handle the complex matrix operations required for density matrix evolution while maintaining numerical stability over extended simulation times.

This framework enables us to track expectation values of arbitrary observables as functions of time, providing direct access to experimentally measurable quantities such as state populations, coherences, and correlation functions. The modular design allows for easy extension to multi-qubit systems and more complex noise models, making it a versatile tool for studying decoherence in realistic quantum systems.

C. Simulation on Quantum Virtual Machine

We conducted simulation of the algorithm we investigate on the Rigetti Quantum Virtual Machine (QVM) using the pyQuil library. The simulation is conducted with error model that are inspired by the Rigetti Ankaa-3 quantum processor and serves as a proof-of-principle demonstration of the QAOA for MAX2SAT problem. Simulation is shown in figure 3

V. EXPERIMENTAL RESULTS AND ANALYSIS

The experiment was conducted on the Rigetti Ankaa-3 quantum processor, the experiment investigates the relation between the ratio of

A. Performance Varying the Clause-to-Variable Ratio

For satisfiability problems, the clause-to-variable ratio $\alpha = m/n$ (where m is the number of clauses and n is the number of variables) serves as a fundamental parameter governing computational difficulty. This ratio captures the constraint density of the problem: low values indicate underconstrained instances with many satisfying




FIG. 2. Temporal evolution of single-qubit state probabilities under four canonical quantum noise channels with identical noise strength $\gamma = 0.1$. Each subplot shows the probability of measuring computational basis state $|0\rangle$ as a function of time. The **bit-flip channel** (upper left) exhibits symmetric oscillatory decay toward the maximally mixed state, modeling random X -gate errors. The **amplitude damping channel** (upper right) demonstrates asymmetric decay with bias toward the ground state, capturing energy dissipation processes. The **phase damping channel** (lower left) preserves population while destroying coherence, representing pure dephasing mechanisms. The **depolarization channel** (lower right) combines all Pauli errors, showing uniform decay toward classical statistics. The oscillatory behavior in all channels reflects the competition between coherent evolution and environmental decoherence, with distinct asymptotic limits revealing the physical nature of each noise mechanism.

assignments, while high values indicate overconstrained instances where satisfying assignments become increasingly rare.

The theoretical significance of this ratio stems from phase transition phenomena in random satisfiability. For MAX-2-SAT, there exists a critical threshold around $\alpha_c \approx 1$ where the probability of satisfiability undergoes a sharp transition [27]. Below this threshold, instances are typically easy to satisfy, while above it, they become increasingly difficult as constraint density grows.

For quantum optimization algorithms like QAOA, the clause-to-variable ratio becomes particularly relevant because quantum advantage may manifest differently across various constraint densities. As α increases, the solution space becomes more constrained, potentially affecting the algorithm's ability to leverage quantum superposition effectively.

Figure 4 demonstrates how QAOA performance varies with constraint density for different problem sizes. The comparison between 3-variable and 5-variable instances reveals how the absolute number of clauses affects performance differently depending on the underlying variable




FIG. 3. QAOA optimization convergence on a quantum virtual machine (QVM) incorporating realistic noise models derived from actual quantum hardware measurements and benchmarking data. This simulation provides a faithful representation of quantum optimization performance under real-world conditions rather than idealized noiseless scenarios. The **quantum expectation** values (blue circles) exhibit characteristic fluctuations that combine both fundamental quantum shot noise and realistic hardware imperfections including gate errors, decoherence, and measurement uncertainties. The **best achieved** score (dark blue line) demonstrates that despite these realistic noise sources, the optimization maintains monotonic progress by tracking the running maximum across all evaluations. The **5-point moving average** (orange dashed line) reveals the underlying convergence trend by filtering out both shot-to-shot statistical variations and correlated noise effects from the hardware model. The optimization exhibits three distinct behavioral regimes: initial exploration with high variance reflecting both algorithmic search and noise amplification (iterations 0-8), accelerated improvement as the algorithm identifies promising parameter regions while noise effects stabilize (iterations 8-15), and final convergence to a near-optimal solution around 4.5 where further improvements become limited by the fundamental noise floor (iterations 15-32). This realistic simulation validates that QAOA can achieve approximately 90% approximation ratios even when subjected to the full spectrum of noise sources present in contemporary quantum processors, providing confidence in the algorithm's viability for near-term quantum advantage demonstrations.

count, highlighting the importance of the normalized α metric.

This systematic variation in performance across different values of α provides crucial insights into where quantum optimization algorithms excel or struggle relative to the underlying problem structure. By using the clause-to-variable ratio as our primary experimental parameter, we can meaningfully compare algorithm performance across different constraint densities while maintaining connection to the rich theoretical framework established in satisfiability research.




FIG. 4. QAOA performance on MAX-2-SAT instances with varying constraint densities. The expected score (normalized as percentage of optimal) is shown for 3-variable instances (blue line with diamonds) and 5-variable instances (orange line with circles) as the number of clauses increases. The vertical dashed line indicates the boundary of the experimental regime. For 3-variable instances, the clause-to-variable ratio ranges from $\alpha = 1$ (3 clauses) to $\alpha = 3$ (9 clauses), while 5-variable instances span from $\alpha = 0.6$ to $\alpha = 1.8$. The performance degradation patterns reveal how constraint density affects algorithm effectiveness across different problem sizes, with 3-variable instances showing more dramatic sensitivity to increased clause counts.

B. Performance Under Equal Variable-Clause Scaling

Figure 5 demonstrates the performance of QAOA when the number of variables and clauses scale together proportionally. This experimental setup maintains a constant constraint density, where each additional variable is accompanied by an additional constraint clause.

The dramatic performance decline from approximately 0.91 to 0.77 as we scale from 3 to 9 variables reveals a fundamental challenge in quantum optimization. When variables and clauses increase together, we observe what complexity theorists call the "constraint satisfaction cliff" - a regime where the solution space becomes increasingly constrained while simultaneously expanding exponentially.

This behavior makes intuitive sense when we consider the underlying quantum mechanics. QAOA relies on quantum superposition to explore multiple solution candidates simultaneously, but as the constraint density remains high with increasing problem size, the quantum state must encode increasingly complex correlations between variables. The gradual performance degradation suggests that the quantum circuit depth and parameter optimization become insufficient to capture these correlations effectively as the problem scales.




FIG. 5. QAOA performance on MAX2SAT instances where the number of variables equals the number of clauses. The expected score represents the fraction of satisfied clauses. Error bars indicate standard deviation across multiple experimental runs. The declining trend reveals how increased problem size with constant constraint density challenges the quantum optimization algorithm.

C. Performance Under Fixed Constraint Scaling

Figure 6 presents a markedly different scaling behavior when the number of clauses remains constant while variables increase. This experimental design reveals how constraint density affects quantum optimization performance.




FIG. 6. QAOA performance on MAX2SAT instances with a fixed number of clauses as variables increase. The initial decline followed by stabilization illustrates how decreasing constraint density eventually benefits quantum optimization, despite the expanding solution space.

The stabilization of performance around 0.8-0.82 after the initial drop from 3 to 4 variables tells a compelling

story about the interplay between solution space size and constraint satisfaction. While adding variables without adding clauses does expand the solution space exponentially, it simultaneously decreases the constraint density - creating more degrees of freedom for the quantum algorithm to find satisfying assignments.

The initial performance drop from 3 to 4 variables likely represents a transition point where the quantum circuit architecture encounters a threshold effect. However, the subsequent stabilization suggests that QAOA can effectively leverage the reduced constraint pressure. Additional unconstrained variables actually provide the quantum algorithm with more flexibility to satisfy the fixed set of clauses, explaining why performance doesn't continue to degrade despite the exponentially growing solution space.

This behavior aligns with our understanding of quantum approximate optimization: the algorithm performs best when it can exploit the structure of under-constrained problems, where quantum superposition can efficiently explore the expanded solution space to find high-quality approximate solutions.

D. Hardware Quality Impact on QAOA Performance

Understanding quantum algorithm performance requires establishing meaningful performance thresholds that distinguish between effective quantum computation and algorithmic failure. In our QAOA implementation for MAX2SAT, we've established prior in this paper 75% as the critical performance boundary below which the algorithm effectively ceases to provide quantum advantage over classical random guessing strategies.

TABLE I. QAOA performance comparison under varying hardware quality conditions for 3 and 4 qubit implementations. Performance values represent approximation ratios as percentages, with 75% marking the theoretical lower bound for meaningful quantum algorithm operation.

Operation Quality	3 Qubits		4 Qubits	
	Theory	Exp	Theory	Exp
High-quality	91.09%	82.40%	89.15%	80.91%
Low-quality	84.58%	76.02%	82.06%	74.00%

Now, here's where the story becomes fascinating from a quantum computing perspective. Notice how the 4-qubit experiment with low-quality operations hits exactly 74.00% - we've reached what you might call the "quantum cliff edge." This isn't just experimental noise or measurement error; we're witnessing the fundamental breakdown of quantum coherence that QAOA requires to function.

Think about what's happening at the quantum level as we cross this threshold. When experimental results drop below 75%, the quantum superposition states that QAOA uses to explore solution spaces simultaneously are

becoming so corrupted by noise that the algorithm can no longer distinguish between good and bad solutions effectively. The quantum interference patterns that create the algorithm's computational advantage are washing out, leaving us with essentially classical random search behavior.

The progression from 3 to 4 qubits reveals something profound about quantum scalability limits under current hardware constraints. While the 3-qubit low-quality experiment maintains 76.02% performance - just above our critical threshold - the 4-qubit case plunges to 74.00%. This tells us we've hit a scaling wall where each additional qubit pushes us closer to the point where quantum noise overwhelms quantum advantage.

Consider why this matters for the broader field of quantum computing. The fact that we can predict and measure these performance boundaries means we're developing a quantitative understanding of where current quantum hardware becomes useful versus where it fails. The 4-qubit low-quality result at 74.00% represents a crucial data point because it shows us the exact conditions under which quantum algorithms transition from providing computational advantage to becoming ineffective.

E. Performance with Number of QAOA Layers

Literature on QAOA suggests the number of layers p remains an open question in the context of MAX2SAT problems. Simulations have been conducted and the result observed is that there exist slight performance improvement in the low clauses number regime, the implementation of such experiment on real quantum hardware was not conducted in this work due to limited hardware resources. The simulation results are shown in the appendix section.

VI. CONCLUSION

This work establishes a comprehensive experimental framework for understanding QAOA performance on MAX2SAT problems, providing both theoretical insights and practical implementation guidelines for near-term quantum optimization. Through systematic experimentation on the Rigetti Ankaa-3 quantum processor, we have identified critical performance boundaries that define when quantum algorithms provide computational advantage over classical alternatives.

A. Research Contributions and Practical Implications

Our circuit construction methodology demonstrates how MAX2SAT problems can be systematically translated into executable quantum circuits through careful decomposition of cost and mixing unitaries. The

single-layer QAOA implementation achieves the theoretically predicted 3/4 approximation ratio, confirming that even shallow quantum circuits can match sophisticated classical-quantum hybrid approaches. This result validates the fundamental promise of quantum approximate optimization while establishing realistic performance expectations for current hardware implementations.

The identification of the clause-to-variable ratio α as a critical performance parameter provides new insights into quantum optimization behavior. Our experiments reveal that QAOA performance exhibits strong dependence on constraint density, with optimal operating regimes occurring near the satisfiability phase transition around $\alpha \approx 1$. This finding connects quantum optimization performance to fundamental satisfiability theory, suggesting that quantum algorithms excel in specific structural regimes rather than providing universal improvements.

The scaling behavior analysis reveals fundamental limitations of current quantum optimization approaches. Equal variable-clause scaling demonstrates how quantum algorithms struggle with increasing constraint complexity, showing performance degradation from 91% to 77% as problems grow. Conversely, fixed constraint scaling exhibits more favorable behavior, with performance stabilizing around 80-82% after initial drops, indicating that reduced constraint density can partially compensate for expanded solution spaces.

Most significantly, our hardware quality experiments establish quantitative thresholds for quantum advantage. The discovery that 75% approximation ratio serves as a critical boundary—below which quantum algorithms lose advantage over classical random strategies—provides a concrete metric for evaluating quantum optimization viability. The experimental demonstration of the “quantum cliff edge” at 74% performance under low-quality 4-qubit conditions illustrates how noise accumulation can push quantum algorithms below this critical threshold. The establishment of the 75% threshold provides researchers and practitioners with a clear target for quantum advantage claims, while the precise measurement of performance under varying hardware conditions enables predictive assessment of when quantum algorithms will succeed or fail on specific hardware platforms.

B. Future Directions and Broader Impact

This work opens several promising research directions that emerge naturally from our experimental findings. The connection between clause-to-variable ratio and quantum performance suggests that adaptive problem decomposition strategies could enhance quantum optimization effectiveness by identifying and exploiting optimal constraint density regimes. The hardware quality analysis framework could be extended to other quantum algorithms and hardware platforms, enabling systematic characterization of quantum advantage boundaries across different computational tasks. The observed scaling behaviors indicate opportunities for hybrid quantum-classical approaches that leverage quantum algorithms in optimal constraint density regimes while employing classical methods for other problem instances.

Beyond immediate quantum computing applications, this work contributes to our fundamental understanding of where quantum algorithms provide computational advantages. By establishing quantitative performance boundaries and connecting them to underlying physical limitations, we advance the theoretical foundation for quantum advantage claims while providing practical tools for algorithm evaluation. The experimental methodology developed here—combining systematic circuit construction, comprehensive performance analysis, and hardware quality assessment—provides a template for rigorous evaluation of quantum algorithms on near-term devices. This approach ensures that quantum advantage claims are grounded in careful experimental validation rather than idealized theoretical projections.

The precise characterization of noise effects on quantum optimization performance provides foundation for developing error mitigation strategies specifically tailored to variational quantum algorithms. Additionally, the quantitative thresholds and scaling relationships established here provide essential guidance for the continued development of practically useful quantum optimization algorithms. While current quantum optimization algorithms face significant scaling and noise limitations, this work demonstrates that meaningful quantum advantages can be achieved within specific problem regimes using careful algorithm design and realistic performance assessment.

-
- [1] P. Benioff, The computer as a physical system: A microscopic quantum mechanical hamiltonian model of computers as represented by turing machines, *Journal of Statistical Physics* **22**, 563 (1980).
 - [2] R. P. Feynman, Simulating physics with computers, *International journal of theoretical physics* **21**, 467 (1982).
 - [3] D. Deutsch, Quantum theory, the church-turing principle and the universal quantum computer, *Proceedings of the Royal Society of London. A. Mathematical and Physical Sciences* **400**, 97 (1985).
 - [4] P. W. Shor, Algorithms for quantum computation: Discrete logarithms and factoring, in *Proceedings 35th Annual Symposium on Foundations of Computer Science* (IEEE, 1994) pp. 124–134.
 - [5] L. K. Grover, A fast quantum mechanical algorithm for database search, in *Proceedings of the twenty-eighth annual ACM symposium on Theory of computing* (ACM, 1996) pp. 212–219.
 - [6] R. Acharya, D. A. Abanin, L. Aghababaie-Beni, I. Aleiner, T. I. Andersen, M. Ansmann, F. Arute,

- K. Arya, A. Asfaw, N. Astrakhantsev, J. Atalaya, R. Babbush, D. Bacon, B. Ballard, J. C. Bardin, J. Bausch, A. Bengtsson, A. Bilmes, S. Blackwell, S. Boixo, G. Bortoli, A. Bourassa, J. Bovaird, L. Brill, M. Broughton, D. A. Browne, B. Buchea, B. B. Buckley, D. A. Buell, T. Burger, B. Burkett, N. Bushnell, A. Cabrera, J. Campero, H.-S. Chang, Y. Chen, Z. Chen, B. Chiaro, D. Chik, C. Chou, J. Claes, A. Y. Cleland, J. Cogan, R. Collins, P. Conner, W. Courtney, A. L. Crook, B. Curtin, S. Das, A. Davies, L. De Lorenzo, D. M. Debroy, S. Demura, M. Devoret, A. Di Paolo, P. Donohoe, I. Drozdov, A. Dunsworth, C. Earle, T. Edlich, A. Eickbusch, A. M. Elbag, M. Elzouka, C. Erickson, L. Faoro, E. Farhi, V. S. Ferreira, L. F. Burgos, E. Forati, A. G. Fowler, B. Foxen, S. Ganjam, G. Garcia, R. Gasca, É. Genois, W. Giang, C. Gidney, D. Gilboa, R. Gosula, A. G. Dau, D. Graumann, A. Greene, J. A. Gross, S. Habegger, J. Hall, M. C. Hamilton, M. Hansen, M. P. Harrigan, S. D. Harrington, F. J. H. Heras, S. Heslin, P. Heu, O. Higgott, G. Hill, J. Hilton, G. Holland, S. Hong, H.-Y. Huang, A. Huff, W. J. Huggins, L. B. Ioffe, S. V. Isakov, J. Iveland, E. Jeffrey, Z. Jiang, C. Jones, S. Jordan, C. Joshi, P. Juhas, D. Kafri, H. Kang, A. H. Karamlou, K. Kechedzhi, J. Kelly, T. Khaira, T. Khattar, M. Khezri, S. Kim, P. V. Klimov, A. R. Klotz, B. Kobrin, P. Kohli, A. N. Korotkov, F. Kostritsa, R. Kothari, B. Kozlovskii, J. M. Kreikebaum, V. D. Kurilovich, N. Lacroix, D. Landhuis, T. Lange-Dei, B. W. Langley, P. Laptev, K.-M. Lau, L. Le Guevel, J. Ledford, J. Lee, K. Lee, Y. D. Lensky, S. Leon, B. J. Lester, W. Y. Li, Y. Li, A. T. Lill, W. Liu, W. P. Livingston, A. Locharla, E. Lucero, D. Lundahl, A. Lunt, S. Madhuk, F. D. Malone, A. Maloney, S. Mandrà, J. Manyika, L. S. Martin, O. Martin, S. Martin, C. Maxfield, J. R. McClean, M. McEwen, S. Meeks, A. Megrant, X. Mi, K. C. Miao, A. Mieszala, R. Molavi, S. Molina, S. Montazeri, A. Morvan, R. Movassagh, W. Mruczkiewicz, O. Naaman, M. Neeley, C. Neill, A. Nersisyan, H. Neven, M. Newman, J. H. Ng, A. Nguyen, M. Nguyen, C.-H. Ni, M. Y. Niu, T. E. O'Brien, W. D. Oliver, A. Opremcak, K. Ottosson, A. Petukhov, A. Pizzuto, J. Platt, R. Potter, O. Pritchard, L. P. Pryadko, C. Quintana, G. Ramachandran, M. J. Reagor, J. Redding, D. M. Rhodes, G. Roberts, E. Rosenberg, E. Rosenfeld, P. Roushan, N. C. Rubin, N. Saei, D. Sank, K. Sankaragomathi, K. J. Satzinger, H. F. Schurkus, C. Schuster, A. W. Senior, M. J. Shearn, A. Shorter, N. Shutty, V. Shvarts, S. Singh, V. Sivak, J. Skrzynny, S. Small, V. Smelyanskiy, W. C. Smith, R. D. Somma, S. Springer, G. Sterling, D. Strain, J. Suchard, A. Szasz, A. Szein, D. Thor, A. Torres, M. M. Torunbalci, A. Vaishnav, J. Vargas, S. Vdovichev, G. Vidal, B. Villalonga, C. V. Heidweiller, S. Waltman, S. X. Wang, B. Ware, K. Weber, T. Weidel, T. White, K. Wong, B. W. K. Woo, C. Xing, Z. J. Yao, P. Yeh, B. Ying, J. Yoo, N. Yosri, G. Young, A. Zalcman, Y. Zhang, N. Zhu, N. Zobrist, G. Q. AI, and Collaborators, Quantum error correction below the surface code threshold, *Nature* **638**, 920 (2025).
- [7] S. Bravyi, A. W. Cross, J. M. Gambetta, D. Maslov, P. Rall, and T. J. Yoder, High-threshold and low-overhead fault-tolerant quantum memory, *Nature* **627**, 778 (2024).
- [8] K. Wang, Z. Lu, C. Zhang, G. Liu, J. Chen, Y. Wang, Y. Wu, S. Xu, X. Zhu, F. Jin, Y. Gao, Z. Tan, Z. Cui, N. Wang, Y. Zou, A. Zhang, T. Li, F. Shen, J. Zhong, Z. Bao, Z. Zhu, Y. Han, Y. He, J. Shen, H. Wang, J.-N. Yang, Z. Song, J. Deng, H. Dong, Z.-Z. Sun, W. Li, Q. Ye, S. Jiang, Y. Ma, P.-X. Shen, P. Zhang, H. Li, Q. Guo, Z. Wang, C. Song, H. Wang, and D.-L. Deng, Demonstration of low-overhead quantum error correction codes (2025), arXiv:2505.09684 [quant-ph].
- [9] J. R. McClean, J. Romero, R. Babbush, and A. Aspuru-Guzik, The theory of variational hybrid quantum-classical algorithms, *New Journal of Physics* **18**, 023023 (2016).
- [10] K. Bharti, A. Cervera-Lierta, T. H. Kyaw, T. Haug, S. Alperin-Lea, A. Anand, M. Degroote, H. Heimonen, J. S. Kottmann, T. Menke, W.-K. Mok, S. Sim, L.-C. Kwek, and A. Aspuru-Guzik, Noisy intermediate-scale quantum algorithms, *Reviews of Modern Physics* **94**, 10.1103/revmodphys.94.015004 (2022).
- [11] M. Cerezo, A. Arrasmith, R. Babbush, S. C. Benjamin, S. Endo, K. Fujii, J. R. McClean, K. Mitarai, X. Yuan, L. Cincio, and P. J. Coles, Variational quantum algorithms, *Nature Reviews Physics* **3**, 625 (2021).
- [12] C. Bravo-Prieto, R. LaRose, M. Cerezo, Y. Subasi, L. Cincio, and P. J. Coles, Variational quantum linear solver, *Quantum* **7**, 1188 (2023).
- [13] Z. Jiang, E. G. Rieffel, and Z. Wang, Near-optimal quantum circuit for grover's unstructured search using a transverse field, *Phys. Rev. A* **95**, 062317 (2017).
- [14] C. Kokail, C. Maier, R. van Bijnen, T. Brydges, M. K. Joshi, P. Jurcevic, C. A. Muschik, P. Silvi, R. Blatt, C. F. Roos, and P. Zoller, Self-verifying variational quantum simulation of lattice models, *Nature* **569**, 355 (2019).
- [15] B. Bauer, S. Bravyi, M. Motta, and G. K.-L. Chan, Quantum algorithms for quantum chemistry and quantum materials science, *Chemical Reviews* **120**, 12685 (2020).
- [16] S. Lloyd, M. Schuld, A. Ijaz, J. Izaac, and N. Killoran, Quantum embeddings for machine learning (2020), arXiv:2001.03622 [quant-ph].
- [17] A. Delgado, J. M. Arrazola, S. Jahangiri, Z. Niu, J. Izaac, C. Roberts, and N. Killoran, Variational quantum algorithm for molecular geometry optimization, *Phys. Rev. A* **104**, 052402 (2021).
- [18] D. Amaro, C. Modica, M. Rosenkranz, M. Fiorentini, M. Benedetti, and M. Lubasch, Filtering variational quantum algorithms for combinatorial optimization, *Quantum Science and Technology* **7**, 015021 (2022).
- [19] E. Farhi, J. Goldstone, and S. Gutmann, A quantum approximate optimization algorithm (2014), arXiv:1411.4028 [quant-ph].
- [20] L. Zhou, S.-T. Wang, S. Choi, H. Pichler, and M. D. Lukin, Quantum approximate optimization algorithm: Performance, mechanism, and implementation on near-term devices, *Physical Review X* **10**, 10.1103/physrevx.10.021067 (2020).
- [21] A. Lucas, Ising formulations of many np problems, *Frontiers in Physics* **2**, 10.3389/fphy.2014.00005 (2014).
- [22] R. Babbush, B. O'Gorman, and A. Aspuru-Guzik, Resource efficient gadgets for compiling adiabatic quantum optimization problems, *Annalen der Physik* **525**, 877–888 (2013).
- [23] D. S. Johnson, Approximation algorithms for combinatorial problems, *Journal of Computer and System Sciences* **9**, 256 (1974).

- [24] N. Alon and J. H. Spencer, *The Probabilistic Method*, 3rd ed. (John Wiley & Sons, Hoboken, NJ, 2008).
- [25] M. X. Goemans and D. P. Williamson, New $\frac{3}{4}$ -approximation algorithms for the maximum satisfiability problem, *SIAM Journal on Discrete Mathematics* **7**, 656 (1994).
- [26] M. J. D. Powell, Direct search algorithms for optimization calculations, *Acta Numerica* **7**, 287 (1998).
- [27] T. Müller, The quantum approximate optimization algorithm-perspectives of analog compilation and incoherent evolution, PhD Thesis, Universität des Saarlandes **1**, 287 (2023).




FIG. 7. QAOA performance on MAX2SAT problem with 1 layer of circuit, the number of clauses is 3/5/7 and the number of variables is 3/5/7 from top to bottom respectively. The expected score represents the fraction of satisfied clauses. Error bars indicate standard deviation across multiple experimental runs. The performance is around 80.4%, 86.2% and 89.9% respectively.

Appendix A: QAOA Layer Performance Simulation

Here present the simulation results for QAOA performance with varying number of layers p . The simulation is conducted on the Rigetti Quantum Virtual Machine (QVM) using the pyQuil library.




FIG. 8. QAOA performance on MAX2SAT problem with 2 layer of circuit, the number of clauses is 3/5/7 and the number of variables is 3/5/7 from top to bottom respectively. The expected score represents the fraction of satisfied clauses. Error bars indicate standard deviation across multiple experimental runs. The performance is around 83.3%, 91.1% and 88.8% respectively.